

# Key-Value Databases

UA.DETI.CBD

José Luis Oliveira / Carlos Costa

# Outline

---

- ❖ Key-value stores
  - General principles
- ❖ Riak, Redis
  - Characteristics
  - Main Features
  - Use Cases

# Key-Value Databases

---

- ❖ Key value stores are the simplest of NOSQL types
  - consisting only of a unique key and a bucket containing any data you wish to store there.
- ❖ Key-value pairs
  - Key (id, identifier, primary key) – usually a string.
  - Value: can be anything (text, structure, image, etc.) – a black box for the database system.
- ❖ The content of the bucket can be literally anything
  - But unstructured or semi-structured data are the most common.
- ❖ The buckets can hold quite large entries including BLOBs (Basic Large Objects).
- ❖ KVs are row based systems designed for efficiency.

# Key-Value Databases – Advantages

---

- ❖ Highly fault tolerant – always available.
- ❖ Schema-less offers easier upgrade path for changing data requirements
  - (Document stores provide even greater flexibility).
- ❖ Efficient at retrieving information about a particular object (bucket) with a minimum of disc operations.
- ❖ Very simple data model. Very fast to set up and deploy.
- ❖ Great at scaling horizontally across hundreds or thousands of servers.

# Key-Value Databases – Advantages

---

- ❖ No requirement for SQL queries, indexes, triggers, stored procedures, temporary tables, forms, views, or the other technical overheads of RDBMS.
- ❖ Very high data ingest rates.
  - Favors write once, read many applications.
- ❖ Powerful offline reporting with very large data sets.
- ❖ Some vendors are offering advanced forms of KVs that approach the capabilities of document stores or column oriented stores.

# Key-Value Databases – Disadvantages

---

- ❖ Not suitable for complex applications.
- ❖ Not efficient at updating records where only a portion of a bucket is to be updated.
- ❖ Not efficient at retrieving limited data from specific records.
  - For example, in an employee database returning only records of employees making between \$40K and \$60K.
- ❖ As the volume of data increases maintaining unique values as keys becomes more difficult
  - Some more complexity in generating character strings that will remain unique over a large set of keys.
- ❖ Generally needs to read all the records in a bucket or you may need to construct secondary indexes.

# Key-Value Databases

---

## ❖ **Suitable** use cases

- Session data, user profiles, user preferences, shopping carts, ...
- Create ever-growing datasets that are rarely accessed but grow over time. (Caching)
- Where write performance is your highest priority.

## ❖ **When not** to use

- Relationships among entities
- Queries requiring access to the content of the value part
- Set operations involving multiple key-value pairs

# Key-Value Databases

---



redis





# Key Management

---

❖ How the keys should actually be designed?

❖ **Manually assigned keys**

- Real-world natural identifiers
- E.g. e-mail addresses, login names, ...

❖ **Automatically generated keys**

- Auto-increment integers
  - Not suitable in peer-to-peer architectures!
- More complex keys generated by algorithms
  - Keys composed from multiple components such as time stamps, cluster node identifiers, ...
  - Used in practice

# Query Patterns

---

## ❖ Basic **CRUD** operations

- Only when a key is provided
- The knowledge of the keys is essential
- It might even be difficult for a particular database system to provide a list of all the available keys!

## ❖ **No searching by value**

- But we could instruct the database how to parse the values
- ... so that we can fetch the intended search criteria
- ... and store the references within index structures

## ❖ **Batch / sequential processing**

- MapReduce

# Other Functionality

---

- ❖ **Expiration** of key-value pairs
  - After a certain interval of time key-value pairs are automatically removed from the database
  - Useful for user sessions, shopping carts etc.
- ❖ **Collections** of values
  - We can store not only ordinary values, but also their collections such as ordered lists, unordered sets etc.
- ❖ **Links** between key-value pairs
  - Values can mutually be interconnected via links
  - These links can be traversed when querying
- ❖ Particular functionality depends on the store.

# Riak Key-Value Store



# RiakKV

---

- ❖ Developed by Basho Technologies
  - <http://basho.com/products/riak-kv/>
  - Implemented in Erlang
  - Initial release in 2009
  - Operating system: Linux, Mac OS X, ... (not Windows)
- ❖ Open source, incremental scalability, high availability, operational simplicity, decentralized design, automatic data distribution, advanced replication, fault tolerance, ...
- ❖ General-purpose, concurrent, garbage-collected programming language and runtime system

# Data Model

---

- ❖ Instance (→ bucket types) → buckets → objects
- ❖ **Bucket** = collection of objects (logical, not physical collection)
  - Each object must have a unique key
  - Various properties are set at the level of buckets
    - E.g. default replication factor, read / write quora, ...
- ❖ **Object** = key-value pair
  - Key is a Unicode string
  - Value can be anything (text, binary object, image, ...)

Each object is also associated with metadata

  - E.g. its content type (text/plain, image/jpeg, ...),
  - and other internal metadata as well

# Data Model

---

- ❖ How buckets, keys and values should be designed?
- ❖ Complex objects containing various kinds of data
  - E.g. one key-value pair holding information about all the actors and movies at the same time
- ❖ Buckets with different kinds of objects
  - E.g. distinct objects for actors and movies, but all in one bucket
  - Structured naming convention for keys might help
    - E.g. actor\_trojan, movie\_medvidek
- ❖ Separate buckets for different kinds of objects
  - E.g. one bucket for actors, one for movies

# Riak Operations

---

## ❖ Basic CRUD operations

- Create: POST or PUT methods
  - Inserts a key-value pair into a given bucket
  - Key is specified manually, or will be generated automatically
- Read: GET method
  - Retrieves a key-value pair from a given bucket
- Update: PUT method
  - Updates a key-value pair in a given bucket
- Delete: DELETE method
  - Removes a key-value pair from a given bucket

## ❖ Extended functionality

- Links – relationships between objects and their traversal
- Search 2.0 – full-text queries accessing values of objects
- MapReduce



# Riak Usage: API

---

## ❖ HTTP API

- All the user requests are submitted as HTTP requests with an appropriately selected method and specifically constructed URL, headers, and data.
- Example
  - GET /types/<type>/buckets/<bucket>/keys/<key>

## ❖ Protocol Buffers API

## ❖ Erlang API

## ❖ Client libraries for a variety of programming languages

- Official: Java, Ruby, Python, C#, PHP, ...
- Community: C, C++, Haskell, Perl, Python, Scala, ...

# Redis

## (REmote DIctionary Service)



# Redis Overview

---

## ❖ Redis

- **In-memory** key-value store
- Open source, master-slave replication architecture, sharding, high availability, various persistence levels, ...

## ❖ Developed by Redis Labs

## ❖ Implemented in C

## ❖ First release in 2009

## ❖ Available at <http://redis.io/>

# Redis Overview

---

## ❖ Functionality

- Standard key-value store
- Support for structured values (e.g. lists, sets, ...)
- Time-to-live
- Transactions

## ❖ Redis is not just a plain key-value store, but a data structures server, supporting different kind of values.

## ❖ Real-world users

- Twitter, GitHub, Pinterest, StackOverflow, Flickr, ...

# Data Model

---

## ❖ Structure

- Instance → databases → objects

## ❖ **Database** = collection of objects

- Databases do not have names, but integer identifiers [0-15]

## ❖ **Object** = key-value pair

- Key is a string (i.e. any binary data)
- Values can be...
  - Atomic: string
  - Structured: list, set, ordered set, hash

# Data Types

---

## ❖ String

- The only atomic data type
- May contain any binary data (e.g. string, integer counter, PNG image, ...)
- Maximal allowed size is 512 MB

## ❖ List

- Ordered collection of strings
- Elements should preferably be read / written at the head / tail

# Data Types

---

## ❖ Set

- Unordered collection of strings
- Duplicate values are not allowed

## ❖ Sorted set

- Ordered collection of strings
- The order is given by a score (floating number value) associated with each element (from the smallest to the greatest score)

## ❖ Hash

- Associative map between string fields and string values
- Field names have to be mutually distinct

# Interface

---

## ❖ Command line client

- redis-cli

## ❖ Two modes are available...

### ❖ Basic

- Commands are passed as standard command line arguments
  - E.g. redis-cli PING
- Batch processing is possible as well
  - E.g. cat script.txt | redis-cli

### ❖ Interactive

- Users type database commands at the prompt redis-cli

## ❖ **RESP** (REdis Serialization Protocol)



# Basic Commands

---

## ❖ **SELECT [0-15]**

- Select a database (default is 0)

## ❖ **SET** key value

- inserts / replaces a given string

## ❖ **GET** key

- returns a given string

## ❖ **MOVE [key] [db]**

- move key to another database

## ❖ **DBSIZE**

## ❖ **HELP** command

- Provides basic information about Redis commands

# Basic Commands

---

## ❖ **FLUSHDB**

- Deletes all the keys of the currently selected database

## ❖ **FLUSHALL**

- delete all the keys in all the databases

## ❖ **SAVE / BGSAVE**

- Saves the current dataset directly / on background

## ❖ **MONITOR**

- what's going on against your redis datastore (check also redis-stat)

# Strings Operations

---

## ❖ **STRLEN** key

- returns a string length

## ❖ **APPEND** key value

- appends a value at the end of a string

## ❖ **GETRANGE** key start end

- returns a substring Both the boundaries are considered to be inclusive
- Positions start at 0;
- Negative offsets for positions starting at the end

## ❖ **SETRANGE** key offset value

- replaces a substring
- Binary 0 are padded when the original string is not long enough

# Counter Operations

---

- ❖ **INCR** key
- ❖ **DECR** key
  - Increments / decrements a value by 1
- ❖ **INCRBY** key increment
- ❖ **DECRBY** key increment
  - Increments / decrements a value by a given amount

# Handling Keys

---

- ❖ **EXISTS** key
  - determines whether a key exists
- ❖ **KEYS** pattern
  - finds all the keys matching a pattern (\*, ?, ...)
  - E.g. KEYS \*
- ❖ **DEL** key ...
  - removes a given object / objects
- ❖ **RENAME** key newkey
  - changes the key of a given object
- ❖ **TYPE** key – determines the type of a given object
  - Types: string, list, set, zset and hash

# Volatile Keys

---

- ❖ Keys with limited time to live
  - When a specified timeout elapses, a given object is removed
  - Works with any data type
- ❖ **EXPIRE** key seconds
  - Sets a timeout for a given object, i.e. makes the object volatile
  - Can be called repeatedly to change the timeout
- ❖ **TTL** key
  - Returns the remaining time to live for a key
- ❖ **PERSIST** key
  - Removes the existing timeout

# Complex Datatypes

---

- ❖ Redis' popularity comes mostly by supporting:
  - lists, hashes, sets, and sorted sets
- ❖ These collection can contain up to  $2^{32}$  elements (more than 4 billion) per key.
- ❖ Commands follow a good pattern.
  - Set commands begin with S,
  - Hashes with H
  - Sorted sets with Z.
  - List commands generally start with either an L (for left) or an R (for right),
    - depending on the direction of the operation (such as LPUSH).

# Lists

---

- ❖ **LPUSH** key value
- ❖ **RPUSH** key value
  - Adds a new element to the head / tail (Left / Right)
- ❖ **LINSERT** key BEFORE | AFTER pivot value
  - Inserts an element before / after another one
- ❖ **LPOP** key
- ❖ **RPOP** key
  - Removes and returns the first / last element (Left / Right)



# Lists

---

## ❖ **INDEX** key index

- gets an element by its index
  - The first item is at position 0;

## ❖ **LRANGE** key start stop

- gets a range of elements

## ❖ **LREM** key count value

- Removes a "count" number of elements equals to value
- count:
  - Positive / negative = moving from head to tail / tail to head
  - 0 = all the items equals to value are removed

## ❖ **LLEN** key

- gets the length of a list

# Sets

---

- ❖ **SADD** key value ...
  - Adds an element / elements into a set
- ❖ **SREM** key value ...
  - Removes an element / elements from a set
- ❖ **SISMEMBER** key value
  - Determines whether a set contains a given element
- ❖ **SMEMBERS** key
  - gets all the elements of a set
- ❖ **SCARD** key
  - gets the number of elements in a set
- ❖ **SUNION / SINTER / SDIFF** key ...
  - Calculates and returns a set union / intersection / difference of two or more sets

# Hashes

---

- ❖ **HSET** key field value
  - sets the value of a hash field
- ❖ **HGET** key field
  - gets the value of a hash field

## Batch alternatives

- ❖ **HMSET** key field value ... ..
  - Sets values of multiple fields of a given hash
- ❖ **HMGET** key field ...
  - Gets values of multiple fields of a given hash

# Hashes

---

- ❖ **HEXISTS** key field
  - determines whether a given field exists
- ❖ **HGETALL** key
  - gets all the fields and values
- ❖ **HKEYS** key
  - gets all the fields in a given hash
- ❖ **HVALS** key
  - gets all the values in a given hash
- ❖ **HDEL** key field
  - Removes a given field / fields from a hash
- ❖ **HLEN** key
  - returns the number of fields in a given hash

# Sorted Sets

---

## Basic operations

- ❖ **ZADD** key score value ... ..
  - Inserts one element / multiple elements into a sorted set
- ❖ **ZREM** key value ...
  - Removes one element / multiple elements from sorted set

## Working with score

- ❖ **ZSCORE** key value
  - Gets the score associated with a given element
- ❖ **ZINCRBY** key increment value
  - Increments the score of a given element

# Sorted Sets

---

## Retrieval of elements

### ❖ **ZRANGE** key start stop

- Returns all the elements within a given range based on positions

### ❖ **ZRANGEBYSCORE** key min max

- Returns the elements within a given range based on scores

## Other operations

### ❖ **ZCARD** key

- Gets the overall number of all elements

### ❖ **ZCOUNT** key min max

- Counts elements within a given range based on score

# Geospatial field operations

---

- ❖ **GEOADD** key longitude latitude member ...
  - Adds the specified geospatial items (latitude, longitude, name) to the specified key.
- ❖ **GEODIST** key member1 member2 ...
  - Return the distance between two members.
- ❖ **GEOHASH** key member ...
  - Return Geohash string (compatible with [geohash.org](http://geohash.org))
- ❖ **GEOPOS** key member ...
  - Return the positions (longitude,latitude) of all the specified members.
- ❖ **GEOADIUS** key longitude latitude radius ...
  - Return the members which are within the radius of the location.

# RDBMS to Redis: Data Modeling

---

## ❖ Employees: Table

employee_id	first_name	last_name	address
1	John	Doe	New York
2	Benjamin	Button	Chicago
3	Mycroft	Holmes	London

- ❖ In general, any RDBMS table can be represented in a key-value schema as follows:

**`$table_name:$primary_key_value:$attribute_name = $value`**



# RDBMS to Redis: Data Modeling

---

employee_id	first_name	last_name	address
1	John	Doe	New York
2	Benjamin	Button	Chicago
3	Mycroft	Holmes	London

employee:1:first\_name = "John"  
employee:1:last\_name = "Doe"  
employee:1:address = "New York"

employee:2:first\_name = "Benjamin"  
employee:2:last\_name = "Button"  
employee:2:address = "Chicago"

employee:3:first\_name = "Mycroft"  
employee:3:last\_name = "Holmes"  
employee:3:address = "London"

# References

---

- ❖ Commands
  - <http://redis.io/commands>
- ❖ Documentation
  - <http://redis.io/documentation>
- ❖ Data types
  - <http://redis.io/topics/data-types>